Quo Vadis bvol

Richard Hartmann
(rh@globalways.net)

Globalways AG

September 22, 2010

Our experience

- We used to run a set of self-baked scripts...
- We migrated to Bacula because it's:
  - Stable
  - Scalable
  - Flexible
  - Fast
- We love it
- Yet, not everything is perfect
- Side notes

Our experience

- We used to run a set of self-baked scripts...
- We migrated to Bacula because it's:
  - Stable
  - Scalable
  - Flexible
  - Fast
- We love it
- Yet, not everything is perfect
- Side notes

Our experience

- We used to run a set of self-baked scripts...
- We migrated to Bacula because it's:
  - Stable
  - Scalable
  - Flexible
  - Fast
- We love it
- Yet, not everything is perfect
- Side notes

Our experience

- We used to run a set of self-baked scripts...
- We migrated to Bacula because it's:
  - Stable
  - Scalable
  - Flexible
  - Fast
- We love it
- Yet, not everything is perfect
- Side notes

Our experience

- We used to run a set of self-baked scripts...
- We migrated to Bacula because it's:
  - Stable
  - Scalable
  - Flexible
  - Fast
- We love it
- Yet, not everything is perfect
- Side notes

Our experience

- We used to run a set of self-baked scripts...
- We migrated to Bacula because it's:
  - Stable
  - Scalable
  - Flexible
  - Fast
- We love it
- Yet, not everything is perfect
- Side notes

Our experience

- We used to run a set of self-baked scripts...
- We migrated to Bacula because it's:
  - Stable
  - Scalable
  - Flexible
  - Fast
- We love it
- Yet, not everything is perfect
- Side notes
  - Still using Bacula 3
  - Each Director and its Storage are a self-contained unit

Our experience

- We used to run a set of self-baked scripts...
- We migrated to Bacula because it's:
  - Stable
  - Scalable
  - Flexible
  - Fast
- We love it
- Yet, not everything is perfect
- Side notes
  - Still using Bacula 3
  - Each Director and its Storage are a self-contained unit

Our experience

- We used to run a set of self-baked scripts...
- We migrated to Bacula because it's:
  - Stable
  - Scalable
  - Flexible
  - Fast
- We love it
- Yet, not everything is perfect
- Side notes
  - Still using Bacula 3
  - Each Director and its Storage are a self-contained unit

Our experience

- We used to run a set of self-baked scripts...
- We migrated to Bacula because it's:
  - Stable
  - Scalable
  - Flexible
  - Fast
- We love it
- Yet, not everything is perfect
- Side notes
  - Still using Bacula 3
  - Each Director and its Storage are a self-contained unit

Our experience

- We used to run a set of self-baked scripts...
- We migrated to Bacula because it's:
  - Stable
  - Scalable
  - Flexible
  - Fast
- We love it
- Yet, not everything is perfect
- Side notes
  - Still using Bacula 3
  - Each Director and its Storage are a self-contained unit

Issues that annoyed us

- Designed to work with tapes
- Manual configuration does not scale
- No review of backup jobs
  - Situations in which no backup was made at all!
  - Not merely an in-house issue; patches sent
- Catalog backups take too long to replay

Issues that annoyed us

- Designed to work with tapes
- Manual configuration does not scale
- No review of backup jobs
  - Situations in which no backup was made at all!
  - Not merely an in-house issue, patches sent
- Catalog backups take too long to replay

Issues that annoyed us

- Designed to work with tapes
- Manual configuration does not scale
- No review of backup jobs
  - Situations in which no backup was made at all!
  - Not merely an in-house issue; patches sent
- Catalog backups take too long to replay

Issues that annoyed us

- Designed to work with tapes
- Manual configuration does not scale
- No review of backup jobs
  - Situations in which no backup was made at all!
  - Not merely an in-house issue; patches sent
- Catalog backups take too long to replay

Issues that annoyed us

- Designed to work with tapes
- Manual configuration does not scale
- No review of backup jobs
  - Situations in which no backup was made at all!
  - Not merely an in-house issue; patches sent
- Catalog backups take too long to replay

Issues that annoyed us

- Designed to work with tapes
- Manual configuration does not scale
- No review of backup jobs
  - Situations in which no backup was made at all!
  - Not merely an in-house issue; patches sent
- Catalog backups take too long to replay

Management/awareness is everything

- In-house host management system
- UUIDs for everything (sample UUID: 1f39f0a9-0172-4273-93fc-2e393fd5d5b4)
- UUIDs are for computers, not people; need for abstraction and automation
- All solutions need to tie in to our management framework
- Aggressive monitoring and verification

Management/awareness is everything

- In-house host management system
- UUIDs for everything (sample UUID: 1f39f0a9-0172-4273-93fc-2e393fd5d5b4)
- UUIDs are for computers, not people; need for abstraction and automation
- All solutions need to tie in to our management framework
- Aggressive monitoring and verification

Management/awareness is everything

- In-house host management system
- UUIDs for everything (sample UUID: 1f39f0a9-0172-4273-93fc-2e393fd5d5b4)
- UUIDs are for computers, not people; need for abstraction and automation
- All solutions need to tie in to our management framework
- Aggressive monitoring and verification

Management/awareness is everything

- In-house host management system
- UUIDs for everything (sample UUID: 1f39f0a9-0172-4273-93fc-2e393fd5d5b4)
- UUIDs are for computers, not people; need for abstraction and automation
- All solutions need to tie in to our management framework
- Aggressive monitoring and verification

Management/awareness is everything

- In-house host management system
- UUIDs for everything (sample UUID: 1f39f0a9-0172-4273-93fc-2e393fd5d5b4)
- UUIDs are for computers, not people; need for abstraction and automation
- All solutions need to tie in to our management framework
- Aggressive monitoring and verification

"One bvol is one tape"

- Bacula has been designed with tapes in mind
  - Not necessarily a downside
  - Still, not suitable for us
- Recycling storage within bvols can spread backups across several bvols
- Depending on the layout, there is no clean separation between customers, hosts, file sets, and backup jobs
- No direct access from our management system

"One bvol is one tape"

- Bacula has been designed with tapes in mind
  - Not necessarily a downside
  - Still, not suitable for us
- Recycling storage within bvols can spread backups across several bvols
- Depending on the layout, there is no clean separation between customers, hosts, file sets, and backup jobs
- No direct access from our management system

"One bvol is one tape"

- Bacula has been designed with tapes in mind
  - Not necessarily a downside
  - Still, not suitable for us
- Recycling storage within bvols can spread backups across several bvols
- Depending on the layout, there is no clean separation between customers, hosts, file sets, and backup jobs
- No direct access from our management system

"One bvol is one tape"

- Bacula has been designed with tapes in mind
  - Not necessarily a downside
  - Still, not suitable for us
- Recycling storage within bvols can spread backups across several bvols
- Depending on the layout, there is no clean separation between customers, hosts, file sets, and backup jobs
- No direct access from our management system

"One bvol is one tape"

- Bacula has been designed with tapes in mind
  - Not necessarily a downside
  - Still, not suitable for us
- Recycling storage within bvols can spread backups across several bvols
- Depending on the layout, there is no clean separation between customers, hosts, file sets, and backup jobs
- No direct access from our management system

"One bvol is one tape"

- Bacula has been designed with tapes in mind
  - Not necessarily a downside
  - Still, not suitable for us
- Recycling storage within bvols can spread backups across several bvols
- Depending on the layout, there is no clean separation between customers, hosts, file sets, and backup jobs
- No direct access from our management system

"One bvol per backup job"

- Solution: Create one bvol per backup job
- Flexible
- Clear directory structure:
  client-id/host-uuid/client-id_host-uuid_fileset_level_date
- Most granular form of backup storage
- Outside verification of:
- Simply delete bvols when they get stale

"One bvol per backup job"

- Solution: Create one bvol per backup job
- Flexible
- Clear directory structure:
  client-id/host-uuid/client-id_host-uuid_fileset_level_date
- Most granular form of backup storage
- Outside verification of:
- Simply delete bvols when they get stale

"One bvol per backup job"

- Solution: Create one bvol per backup job
- Flexible
- Clear directory structure:
  client-id/host-uuid/client-id_host-uuid_fileset_level_date
- Most granular form of backup storage
- Outside verification of:
  - existence
  - expected file sizes
  - quotas per customer, project, or host
- Simply delete bvols when they get stale

"One bvol per backup job"

- Solution: Create one bvol per backup job
- Flexible
- Clear directory structure:
  client-id/host-uuid/client-id_host-uuid_fileset_level_date
- Most granular form of backup storage
- Outside verification of:
  - existence
  - expected file sizes
  - quotas per customer, project, or host
- Simply delete bvols when they get stale

"One bvol per backup job"

- Solution: Create one bvol per backup job
- Flexible
- Clear directory structure:
  client-id/host-uuid/client-id_host-uuid_fileset_level_date
- Most granular form of backup storage
- Outside verification of:
  - existence
  - expected file sizes
  - quotas per customer, project, or host
- Simply delete bvols when they get stale

"One bvol per backup job"

- Solution: Create one bvol per backup job
- Flexible
- Clear directory structure:
  client-id/host-uuid/client-id_host-uuid_fileset_level_date
- Most granular form of backup storage
- Outside verification of:
  - existence
  - expected file sizes
  - quotas per customer, project, or host
- Simply delete bvols when they get stale

"One bvol per backup job"

- Solution: Create one bvol per backup job
- Flexible
- Clear directory structure:
  client-id/host-uuid/client-id_host-uuid_fileset_level_date
- Most granular form of backup storage
- Outside verification of:
  - existence
  - expected file sizes
  - quotas per customer, project, or host
- Simply delete bvols when they get stale

"One bvol per backup job"

- Solution: Create one bvol per backup job
- Flexible
- Clear directory structure:
  client-id/host-uuid/client-id_host-uuid_fileset_level_date
- Most granular form of backup storage
- Outside verification of:
  - existence
  - expected file sizes
  - quotas per customer, project, or host
- Simply delete bvols when they get stale

"One bvol per backup job"

- Solution: Create one bvol per backup job
- Flexible
- Clear directory structure:
  client-id/host-uuid/client-id_host-uuid_fileset_level_date
- Most granular form of backup storage
- Outside verification of:
  - existence
  - expected file sizes
  - quotas per customer, project, or host
- Simply delete bvols when they get stale

We are not Sisyphus; neither should you

- Manual configuration does not scale
  - Takes time
  - Error-prone
  - Hard to verify programmatically
  - We are using UUIDs, not names
- Database-driven config generation

We are not Sisyphus; neither should you

- Manual configuration does not scale
  - Takes time
  - Error-prone
  - Hard to verify programmatically
  - We are using UUIDs, not names
- Database-driven config generation

We are not Sisyphus; neither should you

- Manual configuration does not scale
  - Takes time
  - Error-prone
  - Hard to verify programmatically
  - We are using UUIDs, not names
- Database-driven config generation

We are not Sisyphus; neither should you

- Manual configuration does not scale
  - Takes time
  - Error-prone
  - Hard to verify programmatically
  - We are using UUIDs, not names
- Database-driven config generation

We are not Sisyphus; neither should you

- Manual configuration does not scale
  - Takes time
  - Error-prone
  - Hard to verify programmatically
  - We are using UUIDs, not names
- Database-driven config generation

We are not Sisyphus; neither should you

- Manual configuration does not scale
  - Takes time
  - Error-prone
  - Hard to verify programmatically
  - We are using UUIDs, not names
- Database-driven config generation

From zero to bacula.conf

- Create host in host manager
- Spawn VM from template
- Retrieve UUID from host, copy & paste into host manager
- Select schedule and fileset for host
- Install custom Bacula package
- The same happens on Storage and Director on demand and with verification
- No user interaction required to set up Bacula!

From zero to bacula.conf

- Create host in host manager
- Spawn VM from template
- Retrieve UUID from host, copy & paste into host manager
- Select schedule and fileset for host
- Install custom Bacula package
- The same happens on Storage and Director on demand and with verification
- No user interaction required to set up Bacula!

From zero to bacula.conf

- Create host in host manager
- Spawn VM from template
- Retrieve UUID from host, copy & paste into host manager
- Select schedule and fileset for host
- Install custom Bacula package
  - Sends authentication and host-UUID via XMLRPC
  - Receives configuration
  - Finishes local installation
- The same happens on Storage and Director on demand and with verification
- No user interaction required to set up Bacula!

From zero to bacula.conf

- Create host in host manager
- Spawn VM from template
- Retrieve UUID from host, copy & paste into host manager
- Select schedule and fileset for host
- Install custom Bacula package
  - Sends authentication and host-UUID via XML RPC
  - Receives configuration
  - Finishes local installation
- The same happens on Storage and Director on demand and with verification
- No user interaction required to set up Bacula!

From zero to bacula.conf

- Create host in host manager
- Spawn VM from template
- Retrieve UUID from host, copy & paste into host manager
- Select schedule and fileset for host
- Install custom Bacula package
  - Sends authentication and host-UUID via XML RPC
  - Receives configuration
  - Finishes local installation
- The same happens on Storage and Director on demand and with verification
- No user interaction required to set up Bacula!

From zero to bacula.conf

- Create host in host manager
- Spawn VM from template
- Retrieve UUID from host, copy & paste into host manager
- Select schedule and fileset for host
- Install custom Bacula package
  - Sends authentication and host-UUID via XML RPC
  - Receives configuration
  - Finishes local installation
- The same happens on Storage and Director on demand and with verification
- No user interaction required to set up Bacula!

From zero to bacula.conf

- Create host in host manager
- Spawn VM from template
- Retrieve UUID from host, copy & paste into host manager
- Select schedule and fileset for host
- Install custom Bacula package
  - Sends authentication and host-UUID via XML RPC
  - Receives configuration
  - Finishes local installation
- The same happens on Storage and Director on demand and with verification
- No user interaction required to set up Bacula!

From zero to bacula.conf

- Create host in host manager
- Spawn VM from template
- Retrieve UUID from host, copy & paste into host manager
- Select schedule and fileset for host
- Install custom Bacula package
  - Sends authentication and host-UUID via XML RPC
  - Receives configuration
  - Finishes local installation
- The same happens on Storage and Director on demand and with verification
- No user interaction required to set up Bacula!

From zero to bacula.conf

- Create host in host manager
- Spawn VM from template
- Retrieve UUID from host, copy & paste into host manager
- Select schedule and fileset for host
- Install custom Bacula package
  - Sends authentication and host-UUID via XML RPC
  - Receives configuration
  - Finishes local installation
- The same happens on Storage and Director on demand and with verification
- No user interaction required to set up Bacula!

From zero to bacula.conf

- Create host in host manager
- Spawn VM from template
- Retrieve UUID from host, copy & paste into host manager
- Select schedule and fileset for host
- Install custom Bacula package
  - Sends authentication and host-UUID via XML RPC
  - Receives configuration
  - Finishes local installation
- The same happens on Storage and Director on demand and with verification
- No user interaction required to set up Bacula!

Backups are live-savers; treat them as such

- Custom MySQL plugin to prohibit remote execution of commands
  - Module failed to load but backup "OK"; patch sent upstream
  - Uncaught bug in our module, thus no backup
- Lesson learned: paranoia is good!

Backups are live-savers; treat them as such

- Custom MySQL plugin to prohibit remote execution of commands
  - Module failed to load but backup "OK"; patch sent upstream
  - Uncaught bug in our module, thus no backup
- Lesson learned: paranoia is good!

Backups are live-savers; treat them as such

- Custom MySQL plugin to prohibit remote execution of commands
  - Module failed to load but backup "OK"; patch sent upstream
  - Uncaught bug in our module, thus no backup
- Lesson learned: paranoia is good!

Backups are live-savers; treat them as such

- Custom MySQL plugin to prohibit remote execution of commands
  - Module failed to load but backup "OK"; patch sent upstream
  - Uncaught bug in our module, thus no backup
- Lesson learned: paranoia is good!

Our additional safe-guards

- We introduced additional checks if
  - expected bvols exist
    - monitoring reacts in five minutes or less when new hosts are configured
  - bvols' sizes fit predictions
  - jobs are stuck in state Created for too long
  - any job is not in state Running, Terminated normally, or Canceled
  - non-fatal errors occured
- All data is fed into our monitoring
- Every single event needs to be looked at, verified, commented and closed manually
- Important backups are stored in three separate, remote locations

Our additional safe-guards

- We introduced additional checks if
  - expected bvols exist
    - monitoring reacts in five minutes or less when new hosts are configured
  - bvols' sizes fit predictions
  - jobs are stuck in state Created for too long
  - any job is not in state Running, Terminated normally, or Canceled
  - non-fatal errors occured
- All data is fed into our monitoring
- Every single event needs to be looked at, verified, commented and closed manually
- Important backups are stored in three separate, remote locations

Our additional safe-guards

- We introduced additional checks if
  - expected bvols exist
    - monitoring reacts in five minutes or less when new hosts are configured
  - bvols' sizes fit predictions
  - jobs are stuck in state Created for too long
  - any job is not in state Running, Terminated normally, or Canceled
  - non-fatal errors occured
- All data is fed into our monitoring
- Every single event needs to be looked at, verified, commented and closed manually
- Important backups are stored in three separate, remote locations

Our additional safe-guards

- We introduced additional checks if
  - expected bvols exist
    - monitoring reacts in five minutes or less when new hosts are configured
  - bvols' sizes fit predictions
  - jobs are stuck in state Created for too long
  - any job is not in state Running, Terminated normally, or Canceled
  - non-fatal errors occured
- All data is fed into our monitoring
- Every single event needs to be looked at, verified, commented and closed manually
- Important backups are stored in three separate, remote locations

Our additional safe-guards

- We introduced additional checks if
  - expected bvols exist
    - monitoring reacts in five minutes or less when new hosts are configured
  - bvols' sizes fit predictions
  - jobs are stuck in state Created for too long
  - any job is not in state Running, Terminated normally, or Canceled
  - non-fatal errors occured
- All data is fed into our monitoring
- Every single event needs to be looked at, verified, commented and closed manually
- Important backups are stored in three separate, remote locations

Our additional safe-guards

- We introduced additional checks if
  - expected bvols exist
    - monitoring reacts in five minutes or less when new hosts are configured
  - bvols' sizes fit predictions
  - jobs are stuck in state Created for too long
  - any job is not in state Running, Terminated normally, or Canceled
  - non-fatal errors occured
- All data is fed into our monitoring
- Every single event needs to be looked at, verified, commented and closed manually
- Important backups are stored in three separate, remote locations

Our additional safe-guards

- We introduced additional checks if
  - expected bvols exist
    - monitoring reacts in five minutes or less when new hosts are configured
  - bvols' sizes fit predictions
  - jobs are stuck in state Created for too long
  - any job is not in state Running, Terminated normally, or Canceled
  - non-fatal errors occured
- All data is fed into our monitoring
- Every single event needs to be looked at, verified, commented and closed manually
- Important backups are stored in three separate, remote locations

Our additional safe-guards

- We introduced additional checks if
  - expected bvols exist
    - monitoring reacts in five minutes or less when new hosts are configured
  - bvols' sizes fit predictions
  - jobs are stuck in state Created for too long
  - any job is not in state Running, Terminated normally, or Canceled
  - non-fatal errors occured
- All data is fed into our monitoring
- Every single event needs to be looked at, verified, commented and closed manually
- Important backups are stored in three separate, remote locations

Our additional safe-guards

- We introduced additional checks if
  - expected bvols exist
    - monitoring reacts in five minutes or less when new hosts are configured
  - bvols' sizes fit predictions
  - jobs are stuck in state Created for too long
  - any job is not in state Running, Terminated normally, or Canceled
  - non-fatal errors occured
- All data is fed into our monitoring
- Every single event needs to be looked at, verified, commented and closed manually
- Important backups are stored in three separate, remote locations

Our additional safe-guards

- We introduced additional checks if
  - expected bvols exist
    - monitoring reacts in five minutes or less when new hosts are configured
  - bvols' sizes fit predictions
  - jobs are stuck in state Created for too long
  - any job is not in state Running, Terminated normally, or Canceled
  - non-fatal errors occured
- All data is fed into our monitoring
- Every single event needs to be looked at, verified, commented and closed manually
- Important backups are stored in three separate, remote locations

Is this even a section?

- MySQL hot-copy
- Locks database!
- Takes 20% more space
- Backup faster by 30%
- Lightning fast restoration

Is this even a section?

- MySQL hot-copy
- Locks database!
- Takes 20% more space
- Backup faster by 30%
- Lightning fast restoration

Is this even a section?

- MySQL hot-copy
- Locks database!
- Takes 20% more space
- Backup faster by 30%
- Lightning fast restoration

Is this even a section?

- MySQL hot-copy
- Locks database!
- Takes 20% more space
- Backup faster by 30%
- Lightning fast restoration

Is this even a section?

- MySQL hot-copy
- Locks database!
- Takes 20% more space
- Backup faster by 30%
- Lightning fast restoration

The End

Thank you for your attention!

Richard Hartmann
Globalways AG
rh@globalways.net

RichiH @ irc://irc.freenode.net/bacula