backy

VM backup beyond Bacula/Bareos

Christian Theune
@theuni
ct@flyingcircus.io

Mea Culpa

And I almost missed it — again

Backup!!11!!

FLYING CIRCUS

- flyingcircus.io
- DevOps as a Service
- custom, mission-critical web applications

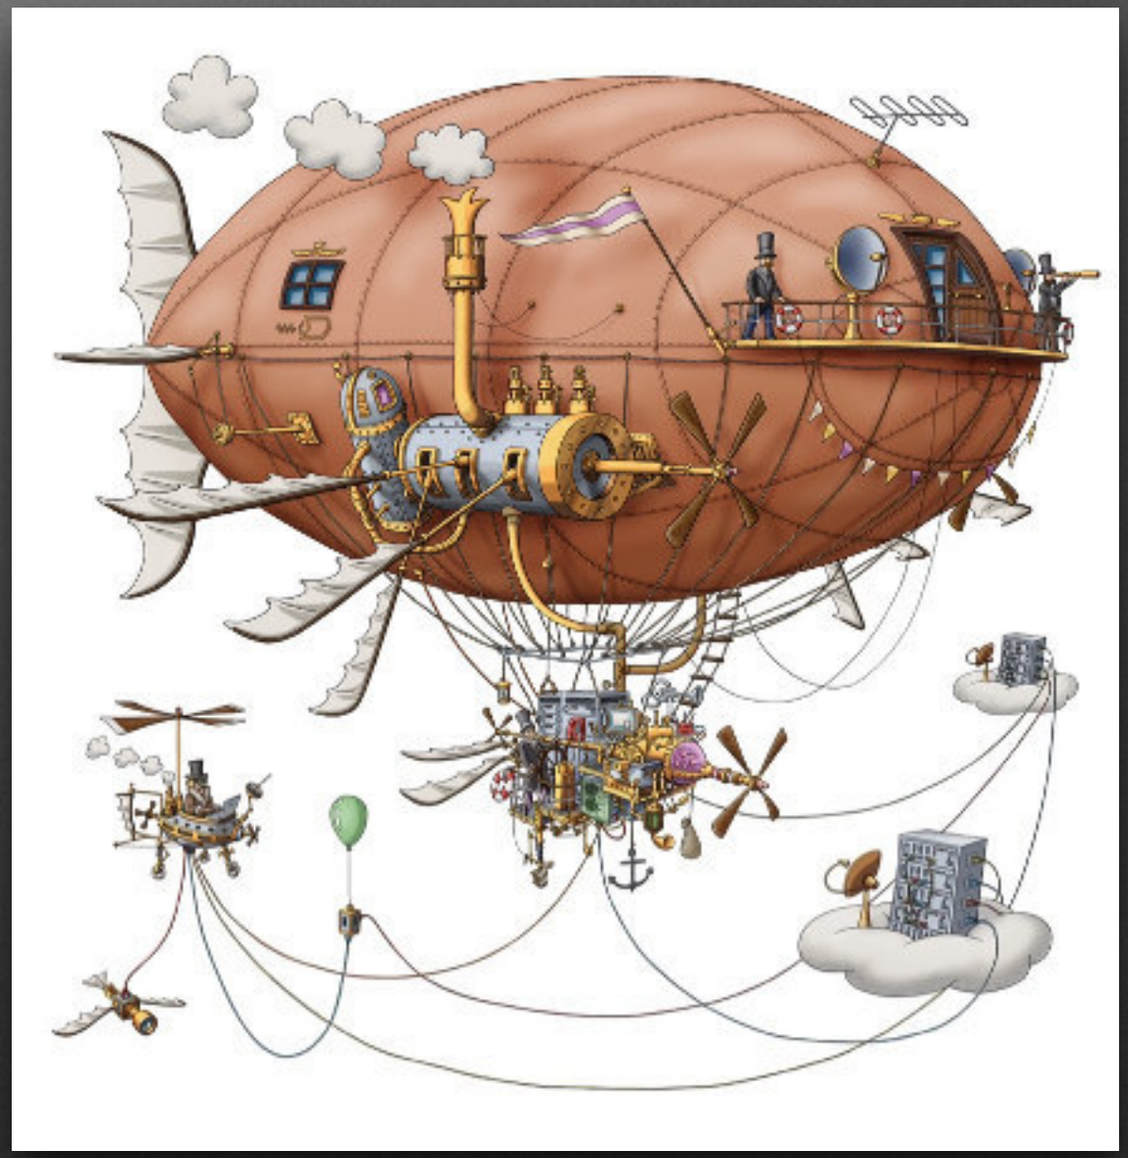

Part I - Oh the Pain

The story unfolds …

Investigating We're seeing disk errors on multiple VMs causing overall service outages. We're currently looking into the issue.

Posted over 1 year ago. Mar 10, 2014 - 14:48 CEST

Identified

We identified the issue of the current outage. We are experiencing a massive data loss after a bug in our management code was triggered when reactivating one of the old storage servers in a foreign location.

Currently we are taking inventory of the damage and preparing for disaster recovery and reinstallation. We'll follow up with a more detailed plan shortly.

Posted over 1 year ago. Mar 10, 2014 - 15:03 CEST

Update

Our restore processes have started and the VMs are recovering. We are initially going to restore central services (mail) and customers with SLA requirements.

We also noticed that not all VMs were affected by the bug - likely due to a timeout causing some of the massive list of deletions to not be executed.

Posted over 1 year ago. Mar 10, 2014 - 16:37 CEST

Update

Almost all production VMs have been restored since around 01:00 CET.

We're currently struggling with a backup inconsistency, having to dig deeper to get a good restore, but overall things are looking OK.

Posted over 1 year ago. Mar 11, 2014 - 05:47 CEST

Resolved All customer VMs have been restored since about 10:00 CET.

Root Cause Analysis

http://flyingcircus.io/postmortems/13266.pdf

Restore script bottleneck: global lock

Undetected inconsistency in important customer database

Bacula: complexity and the VTL

Not “everything” backed up.

24 hours are not a sufficient RPO in quite a few cases

Paper cuts

- Hard link farms
- Boot loaders
- The director as a “most valuable bottleneck”

Recap

- Restore fiddly to script
- Undetected inconsistency that was hard to deal with
- Blind spots
- Daily Interval
- Overall complexity, performance and the VTL
- Paper cuts

Part II - Make a wish

Simplicity

- Restore with basic Unix tools
- No VTL
- Not mixing data of different VMs

Reliability

- Verification / Scrubbing / (Repair)
- High frequency
- Integration with storage snapshots
- Not inventing new formats

Operability

- Avoid bottlenecks / head-of-line blocking
- Efficient deltas for large files (ZODB)
- Parallelisation (multiple jobs and multiple servers)
- Simple scripting and environment-specific integration
- Coordination: pre/post actions on storage, hypervisor, VM …

Operability II

- Simple Nagios integration to ensure we notice RPO/SLA failures
- RTO-compliance during mass-restore
- Self-service for customers to restore files or VMs

Part III - Let’s do this!

It’s all about size

“One size fits all … not”

–Probably someone, maybe me

It’s all about size: backy

```
SLOC    Directory       SLOC-by-Language (Sorted)
1141    tests           python=1141
1048    top_dir         python=1048
857     sources         python=857
0       __pycache__     (none)


Totals grouped by language (dominant language first):
python:          3046 (100.00%)
```

It's all about size: Bacula

SLOC	Directory	SLOC-by-Language (Sorted)
28027	src_lib	ansic=28027
27140	src_dird	ansic=27140
24220	src_stored	ansic=24220
14498	src_qt-console	cpp=14313,sh=185
13015	src_filed	ansic=13015
11289	autoconf	sh=11185,ansic=85,python=11,awk=8
10351	src_cats	ansic=8289,sh=1967,perl=95
5220	src_tools	ansic=5220
4800	examples	perl=2801,sh=1586,ansic=413
3160	src_findlib	ansic=3160
3075	updatedb	sh=3075
2612	src_console	ansic=2612
2376	platforms	sh=1936,ansic=408,python=32
1788	src_plugins	ansic=1788
1371	src_top_dir	cpp=878,ansic=493
1327	scripts	sh=1266,perl=61
403	release	sh=354,perl=49
16	po	sed=16
2	top_dir	sh=2
0	manpages	(none)

```
Totals grouped by language (dom
ansic:          114870 (74.26%)
sh:              21556 (13.93%)
cpp:             15191 (9.82%)
perl:             3006 (1.94%)
python:             43 (0.03%)
sed:                16 (0.01%)
awk:                 8 (0.01%)
```

It’s all about size: Bareos

```
SLOC    Directory       SLOC-by-Language (Sorted)
36681   src_dird        ansic=36681
36643   src_ndmp        ansic=36643
34841   src_lib         ansic=34841
27613   src_stored      ansic=27613
14210   src_cats        ansic=12584,sh=810,lisp=721,perl=95
12797   src_qt-console  cpp=12784,sh=13
12040   src_plugins     ansic=10689,python=1351
11246   autoconf        sh=11227,python=11,awk=8
10475   src_filed       ansic=10475
9225    src_win32       ansic=9063,sh=162
8821    src_findlib     ansic=8821
7959    src_lmdb        ansic=7959
2681    src_console     ansic=2681
2459    src_tools       ansic=2459
2150    platforms       sh=1523,ansic=408,python=189,sed=30
2005    scripts         sh=1944,perl=61
1668    src_tests       ansic=1668
1583    src_qt-tray-monitor cpp=1583
1558    src_include     cpp=1020,ansic=538
953     debian          sh=953
32      test            sh=32
16      po              sed=16
1       top_dir         sh=1
0       manpages        (none)
0       src_defaultconfigs (none)
0       src_images      (none)
0       src_top_dir     (none)
```

```
Totals grouped by language (dom
ansic:          203123 (85.47%)
sh:              16665 (7.01%)
cpp:             15387 (6.47%)
python:           1551 (0.65%)
lisp:              721 (0.30%)
perl:              156 (0.07%)
sed:                46 (0.02%)
awk:                 8 (0.00%)
```